

A TALE OF CONCURRENCY
THROUGH CREATIVITY IN PYTHON:
A DEEP DIVE INTO HOW GEVENT WORKS

KAVYA

What is asynchronous I/O?

What is gevent?

network



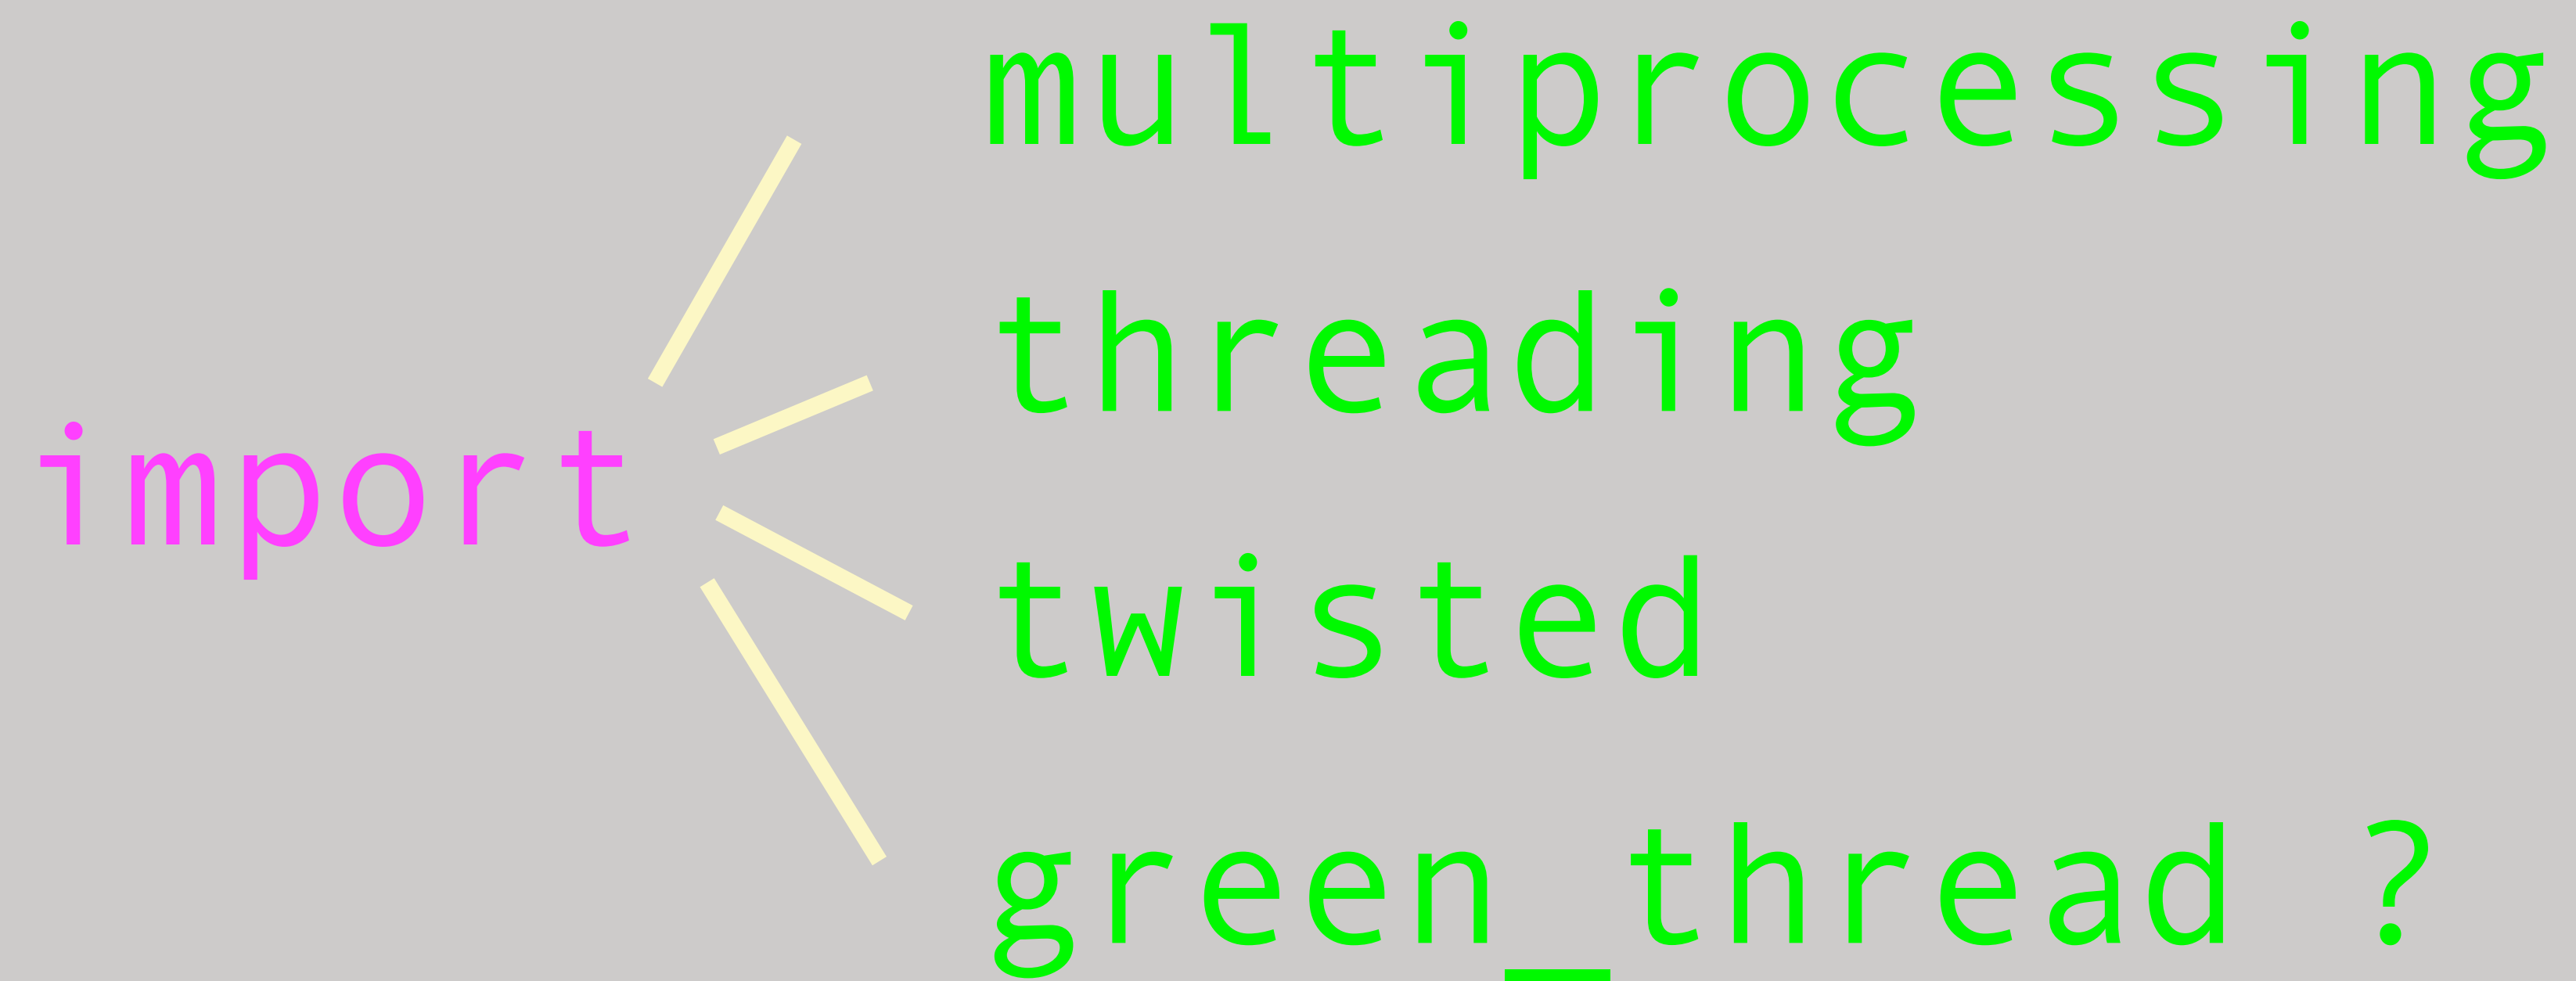
download_photos

```
def download_photos(user):  
    # Open a connection to the server  
    conn = get_authenticated_connection(user)  
  
    # Download all photos  
    photos = get_photos(conn)  
  
    # Save for later display  
    save_photos(user, photos)
```

```
def downloader():  
    users = get_users()  
    for user in users:  
        download_photos(user)
```



network I/O




```
import multiprocessing
```

```
import multiprocessing as mp

def downloader():
    pool = []
    for user in users:
        p = mp.Process(download_photos, user)
        pool.append(p)
        p.start()

    for p in pool:
        p.join()
```

```
import threading
```

```
import threading

def downloader():
    pool = []
    for user in users:
        t = threading.Thread(download_photos, user)
        pool.append(t)
        t.start()

    for t in pool:
        t.join()
```

```
import twisted
```

```
import twisted

def download_photos():
    # Modify this to add callbacks

def downloader():
    # Something something loop.run()
```

green threads

- user space –
the OS does not create or manage them
- cooperatively scheduled –
the OS does not schedule or preempt them
- lightweight

```
import gevent
```



```
import gevent
from gevent import monkey; monkey.patch_all()

def downloader():
    pool = []
    for user in users:
        g = gevent.Greenlet(download_photos,
                               user)

        g.start()
        pool.append(g)
    gevent.joinall(pool)
```

A red arrow originates from the upper right and points diagonally down and to the left, ending its arrowhead directly under the text 'gevent.Greenlet' in the line 'g = gevent.Greenlet(download_photos, user)'. The text 'gevent.Greenlet' is underlined with a red line.

THE BUILDING BLOCKS

PUTTING IT TOGETHER

WRAP-UP/ Q&A

THE BUILDING BLOCKS

```
g = eventloop.run_in_executor(download_photos, user)
```

```
...
```

```
class Greenlet(greenlet):
```

```
    """
```

```
    A light-weight cooperatively-scheduled  
    execution unit.
```

```
    """
```

```
...
```

```
from greenlet import greenlet red
```

```
gr1 = greenlet(print_red)
```

```
gr2 = greenlet(print_blue)
```

```
gr1.switch()
```

blue

red done!

```
def print_red():
```

```
    print 'red' ↓
```

```
    gr2.switch()
```


```
    print 'red done!'
```

```
def print_blue():
```

```
    print 'blue' ↓
```

```
    gr1.switch()
```

```
    print 'blue done!'
```




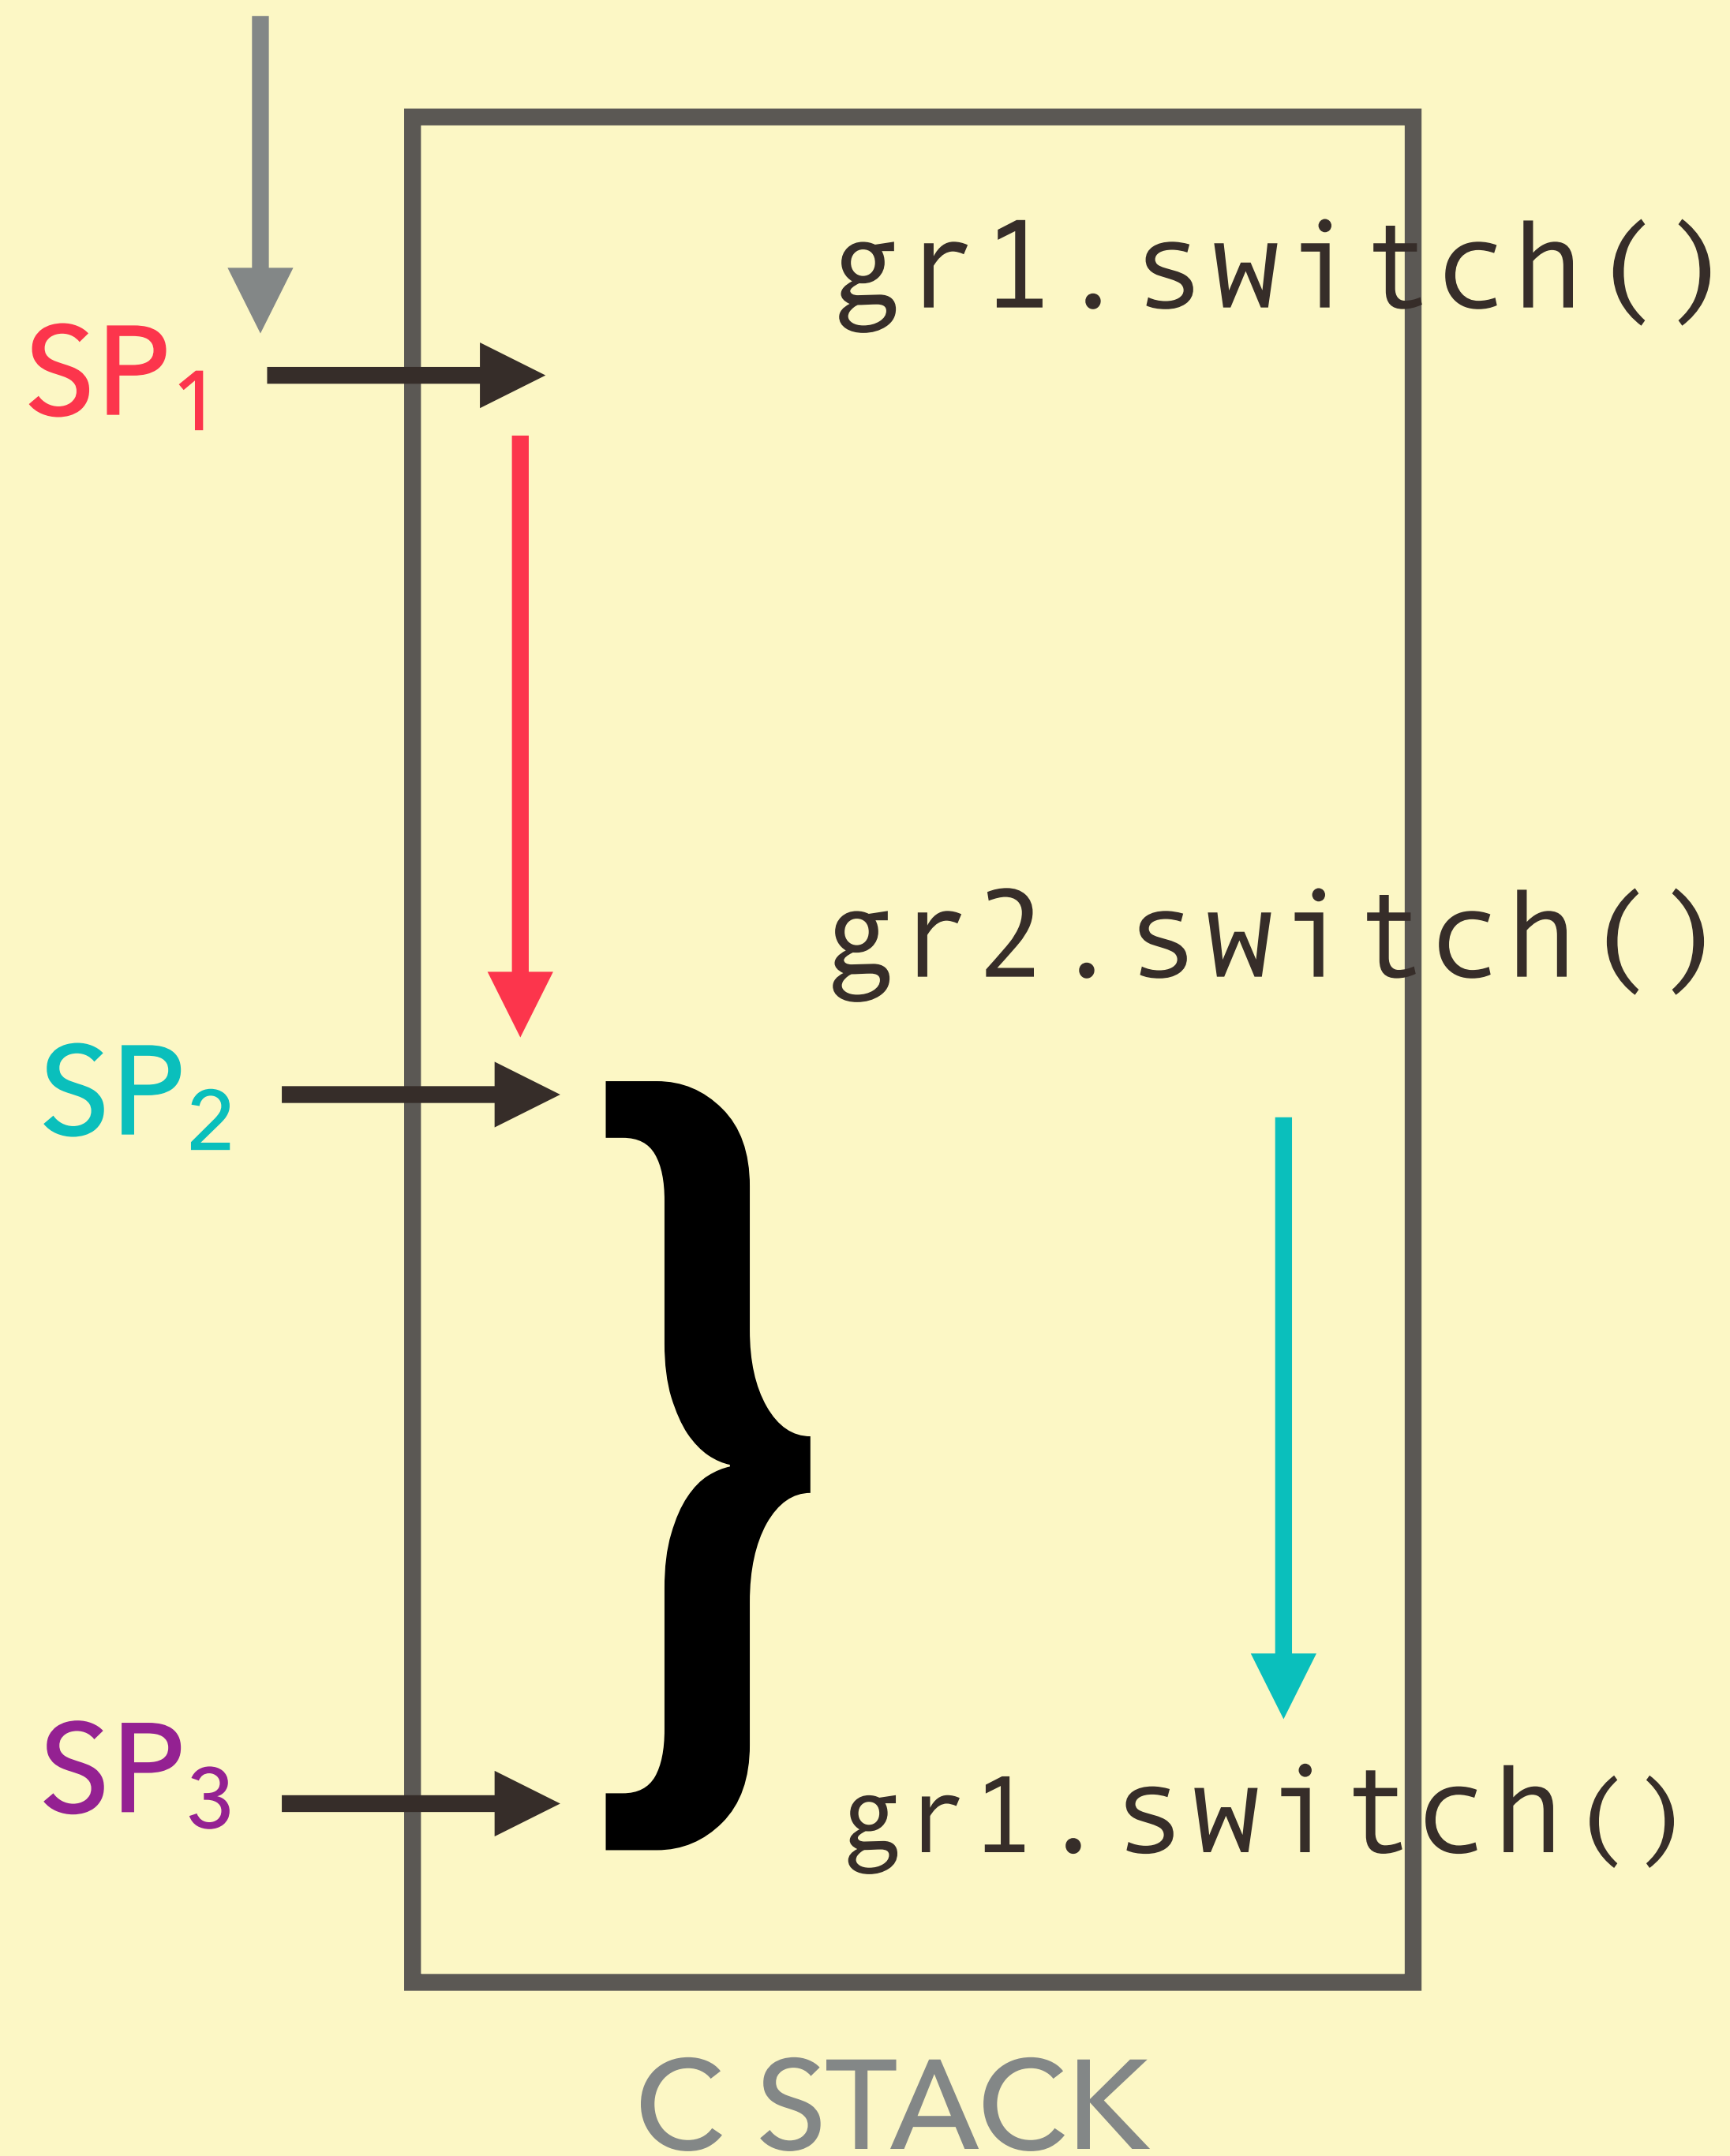
`.switch()`

- pause current + yield control flow
- resume next.`switch()`



`coroutine`


`gr1 = greenlet(run_fn)`  `{`
`run_fn`
`parent`
`...`
`}`



```
{  
    base =  $SP_1$   
}
```

```
{  
    base =  $SP_1$   
    start =  $SP_2$   
}
```

```
{  
    base =  $SP_2$   
}
```

greenlets

for

coroutines

via

assembly-based stack-slicing

```
import gevent
from gevent import monkey; monkey.patch_all()

def downloader():
    pool = []
    for user in users:
        g = gevent.Greenlet(download_photos,
                               user)

        g.start()
        pool.append(g)
    gevent.joinall(pool)
```

`g.start()`

```
def start(self):  
    """ Schedule the greenlet to run in this  
    loop iteration. """  
    if self._start_event is None:  
        self._start_event = \  
            ...loop.run_callback(self.switch)
```

libev

- API to register event_handler callbacks
- watches for events
- calls registered callbacks

“Hey loop,
Wait for a write on this socket and
call `parse_recv()` when that happens.”

```
fd = make_nonblocking(socket_fd)
loop.io_watch(fd, write, callback_fn)
loop.run()
```




```
while True:
```

```
    call all pre_block_watchers
```

```
    block for I/O
```

```
    call pending io_watchers
```

```
    call all post_block_watchers
```



`always` call `pre_block_watchers`

Hook to integrate other event mechanisms
into the loop.



“Hey loop,

If there are coroutines ready to run,
run them. Then, block for a write on...”

libev
for an
event loop

PUTTING IT TOGETHER

```
import gevent
from gevent import monkey; monkey.patch_all()

def downloader():
    pool = []
    for user in users:
        g = gevent.Greenlet(download_photos,
                               user)

        g.start()
        pool.append(g)
    gevent.joinall(pool)
```

```
for user in users:  
    g = gevent.Greenlet(download_photos, user)
```

```
g = gevent.Greenlet(download_photos, user)
```

```
class Greenlet(greenlet):
```

```
    def __init__(self, run=None, ...):
```

```
        greenlet.__init__(self, None, get_hub())
```




```
g.parent = Hub
```


```
class Greenlet(greenlet):  
    greenlet.__init__(self, None, get_hub())
```

↓
`g.parent = Hub`

```
class Hub(greenlet):  
    def __init__(self):  
        greenlet.__init__(self)  
        self.loop = ...
```



Greenlet()



a greenlet –
to run `download_photos()`



the event loop –
i.e. the Hub

`.parent`



```
for user in users:  
    g = gevent.Greenlet(download_photos,user)  
    g.start()
```



```
g.start()
```

Hub

```
self.parent.loop.run_callback(self.switch)
```



```
loop.run()
```

```
while True:
```

```
    call all pre_block_watchers = g.switch
```

```
    block for I/O
```

```
    ...
```

“Hey loop,
.`start()` — This coroutine is `ready to run`.
Run it before you block...”

```
for user in users:  
    g = gevent.Greenlet(download_photos,user)  
    g.start()  
    pool.append(g)  
gevent.joinall(pool)
```

```
gevent.joinall()
```

Hub

```
result = self.parent.switch()
```

```
class Hub(greenlet):
```

```
    def run(self):
```

```
        while True:
```

```
            self.loop.run()
```

`.join()` — runs the loop


```
.join()
```

```
= loop.run() — while True:
```


```
    call pre_block_watchers
```

```
    ... = g.switch
```

```
                |  
                download_photos()
```



HUB



`loop.run()` — `g.switch()`



`download_photos()`



`network I/O`

```
import gevent
```

```
from gevent import monkey; monkey.patch_all()
```

```
def downloader():
```

```
    ...
```

gevent.socket


import ~~socket~~

```
create:  fd = make_nonblocking(socket_fd)
send:    loop.io_watch(fd, write, callback_fn)
        loop.run()
```


Hub.switch

g.switch

network I/O



```
for user in users:  
    g = gevent.Greenlet(download_photos, user)  
    g.start()
```



```
pre_block_watchers = [g1.switch,  
                       g2.switch]
```

```
gevent.joinall()
```

```
Hub loop.run()
```

```
    call_pre_block_watchers = [g1.switch, ...]
```

```
g1.switch()
```

```
g1 download_photos(user1)
```

```
    network_request io_watchers = [g1.switch]
```

```
    Hub.switch()
```

Hub

```
loop.run()
```

```
call_pre_block_watchers = [g2.switch]
```

```
g2.switch()
```

g2

```
download_photos(user2)
```

```
network_request io_watchers = [g2.switch,  
                                g1.switch]
```

```
Hub.switch()
```


Hub

```
loop.run()
```

```
    call pre_block_watchers = []
```

```
    block for I/O
```

```
    call pending io_watchers = [g1.switch]
```

```
g1.switch()
```

g1

```
resumes download_photos(user1)
```

...

WRAP-UP

minuses

- no parallelism
- non-cooperative code will block the entire process:
 - C-extensions → use pure Python libraries
 - compute-bound greenlets → use `gevent.sleep(0)`
 - use greenlet blocking detection
- monkey-patch may have confusing implications
 - order of imports matters

...but


- excellent for workloads that are:
I/O bound, highly concurrent → **20-30k concurrent connections!**
- Used at “web scale” at:
Pinterest, Facebook, Mixpanel, PayPal, Disqus, Nylas...

greenlet

libev

Hub

monkeypatching



greenlet

libev

Hub

monkeypatching