

Transactions Across Datacenters

(and other weekend projects)

Ryan Barrett

May 27, 2009

Of three properties of distributed data systems – consistency, availability, partition tolerance – choose two.

Eric Brewer, CAP theorem, PODC 2000

Scaling is hard.

Various

Context

- *multihoming* (n): operating from multiple datacenters simultaneously
- ~~serving~~
- ~~data processing~~
- ~~read-only data~~
- read/write data (the hardest kind!)
- Case study: App Engine datastore

What's ahead and takeaways

- Consistency?
- Why transactions?
- Why across datacenters?
- Multihoming
- Techniques and tradeoffs
- Conclusions

Consistency

- Weak
- Eventual
- Strong
- Use case: read after write

Weak consistency

- After a write, reads *may or may not* see it
- Best effort only
- "Message in a bottle"
- App Engine: memcache
- VoIP, live online video
- Realtime multiplayer games

Eventual consistency

- After a write, reads *will eventually* see it
- App Engine: mail
- Search engine indexing
- DNS, SMTP, snail mail
- Amazon S3, SimpleDB

Strong consistency

- After a write, reads *will* see it
- App Engine: datastore
- File systems
- RDBMSes
- Azure tables

What's ahead

- Consistency?
- **Why transactions?**
- Why across datacenters?
- Multihoming
- Techniques and tradeoffs
- Conclusions

Why transactions?

- Clichéd example: transfer money from A to B
 - subtract from A
 - add to B
- What if something happens in between?
 - another transaction on A or B
 - machine crashes
 - ...

Why transactions?

- Correctness
- Consistency
- Enforce invariants
- ACID

What's ahead

- Consistency?
- Why transactions?
- **Why across datacenters?**
- Multihoming
- Techniques and tradeoffs
- Conclusions

Why across datacenters?

- Catastrophic failures
- Expected failures
- Routine maintenance
- Geolocality
 - CDNs, edge caching

Why *not* across datacenters?

- Within a datacenter
 - High bandwidth: 1–100Gbps interconnects
 - Low latency: < 1ms within rack, 1–5ms across
 - Little to no cost
- Between datacenters
 - Low bandwidth: 10Mbps–1Gbps
 - High latency: 50–150ms
 - \$\$\$ for fiber

What's ahead

- Consistency?
- Why transactions?
- Why across datacenters?
- **Multihoming**
- Techniques and tradeoffs
- Conclusions

Multihoming

- Hard problem.
- ...consistently? Harder.
- ...with real time writes? Hardest.
- What to do?

Option 1: Don't

- ...instead, bunkerize.
 - Most common
 - Microsoft Azure (tables)
 - Amazon SimpleDB
- Bad at catastrophic failure
 - Large scale data loss
 - Example: SVColo
- Not great for serving
 - No geolocation

Option 2: Primary with hot failover(s)

- Better, but not ideal
 - Mediocre at catastrophic failure
 - Window of lost data
 - Failover data may be inconsistent
- Amazon Web Services
 - EC2, S3, SQS: choose US or EU
 - EC2: Availability Zones, Elastic Load Balancing
- Banks, brokerages, etc.
- Geolocated for reads, not for writes

Option 3: True multihoming

- Simultaneous writes in different DCs
- Two way: hard
- N way: harder
- Handle catastrophic failure, geolocality
- ...but pay for it in latency

What's ahead

- Consistency?
- Why transactions?
- Why across datacenters?
- Multihoming
- **Techniques and tradeoffs**
- Conclusions

Techniques and tradeoffs

	Backups	M/S	MM	2PC	Paxos
Consistency					
Transactions					
Latency					
Throughput					
Data loss					
Failover					

Backups

- Make a copy
- Sledgehammer
- Weak consistency
- Usually no transactions
- Datastore: early internal launch

Backups

	Backups	M/S	MM	2PC	Paxos
Consistency	Weak				
Transactions	No				
Latency	Low				
Throughput	High				
Data loss	Lots				
Failover	Down				

Master/slave replication

- Usually asynchronous
 - Good for throughput, latency
- Most RDBMSes
 - e.g. MySQL binary logs
- Weak/eventual consistency
 - Granularity matters!
- Datastore: current

Master/slave replication

	Backups	M/S	MM	2PC	Paxos
Consistency	Weak	Eventual			
Transactions	No	Full			
Latency	Low				
Throughput	High				
Data loss	Lots	Some			
Failover	Down	Read only			

Multi-master replication

- Umbrella term for merging concurrent writes
- Asynchronous, eventual consistency
- Need *serialization* protocol
 - e.g. *timestamp oracle*: monotonically increasing timestamps
 - Either SPOF with master election...
 - ...or distributed consensus protocol
- No global transactions!
- Datastore: no strong consistency

Multi-master replication

	Backups	M/S	MM	2PC	Paxos
Consistency	Weak	Eventual			
Transactions	No	Full	Local		
Latency	Low				
Throughput	High				
Data loss	Lots	Some			
Failover	Down	Read only	Read/write		

Two Phase Commit

- Semi-distributed consensus protocol
 - deterministic coordinator
- 1: propose, 2: vote, 3: commit/abort
- Heavyweight, synchronous, high latency
- 3PC buys async with extra round trip
- Datastore: poor throughput

Two Phase Commit

	Backups	M/S	MM	2PC	Paxos
Consistency	Weak	Eventual		Strong	
Transactions	No	Full	Local	Full	
Latency	Low			High	
Throughput	High			Low	
Data loss	Lots	Some		None	
Failover	Down	Read only	Read/write		

Paxos

- Fully distributed consensus protocol
- "Either Paxos, or Paxos with craft, or broken"
 - Mike Burrows
- Majority writes; survives minority failure
- Protocol similar to 2PC/3PC
 - Lighter, but still high latency

Paxos for the Datastore

- Close together? no.
- In the *same* datacenter? no.
- Opt-in? maybe...

Paxos for App Engine

- Coordinate moving state
- ...usually via lock server
- Apps
- Memcache
- Offline processing

Paxos

	Backups	M/S	MM	2PC	Paxos
Consistency	Weak	Eventual		Strong	
Transactions	No	Full	Local	Full	
Latency	Low			High	
Throughput	High			Low	Medium
Data loss	Lots	Some		None	
Failover	Down	Read only	Read/write		

Conclusions

- No silver bullet...
- ...but still worth doing!
- Embrace tradeoffs...
- ...and punt them to application developers!

What's behind (pewh!)

- Consistency?
- Why transactions?
- Why across datacenters?
- Multihoming
- Techniques and tradeoffs
- Conclusions

Questions?